

Data structures and Algorithms

B-TREES

Nguyễn Ngọc Thảo
nnthao@fit.hcmus.edu.vn


m-way Trees

The needs of m -way trees

- **Binary trees** allow for **efficient manipulation in memory** due to its simple representation.
 - Each node has **one data item** and **at most two branches**.
- However, it **could not satisfy several practical needs**.
 - Store data neatly in external storage → each unit to be stored should contain more than one data item.
 - Reduce cost for basic operations (i.e., search, insertion or removal) → a better tree architecture is required.
 - Optimal search for a data time → balanced search tree
- **m -way trees are solution to the above issues.**

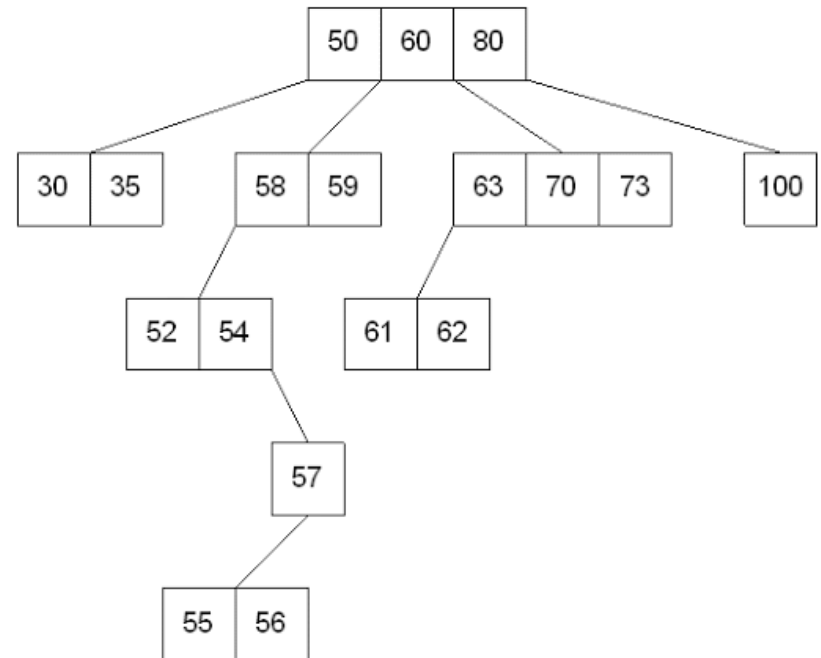
m -way trees: A definition

- A m -way tree has every internal node of at most m children and at most $(m - 1)$ data items.
- The data items within a node are sorted in ascending order.
 - For any i^{th} data item, it is larger than every item in the i^{th} subtree, and smaller than every item in the $(i + 1)^{th}$ subtree.




Example: An example of m -way tree

4-way tree



5-way tree





Insertion in a m -way tree

- Let v be the item to be inserted into a m -way tree.
- Traverse the tree until an empty subtree is found
- If the parent still has empty slot(s), insert v .
- Otherwise, create a new node and insert v into that node.



Example: Insert data item to a 3-way tree

Insert data item 8: Still an empty slot at node $\langle 6 \rangle$.



Example: Insert data item to a 3-way tree

Insert data item 27: No empty slot at node $\langle 28, 30 \rangle$.





Removal in a m -way tree

- Let v be the item to be remove from a m -way tree.
- If v has no child (i.e., it is in between two empty subtrees) then delete v .
- Otherwise, find a substitution for v which is either an in-order predecessor or an in-order successor.



Example: Remove data item from a 3-way tree

Remove data item 8: Node $\langle 6, 8 \rangle$ does not have the 2nd child and 3th child.



Example: Remove data item from a 3-way tree

Remove data item 16: The 1st subtree is available.



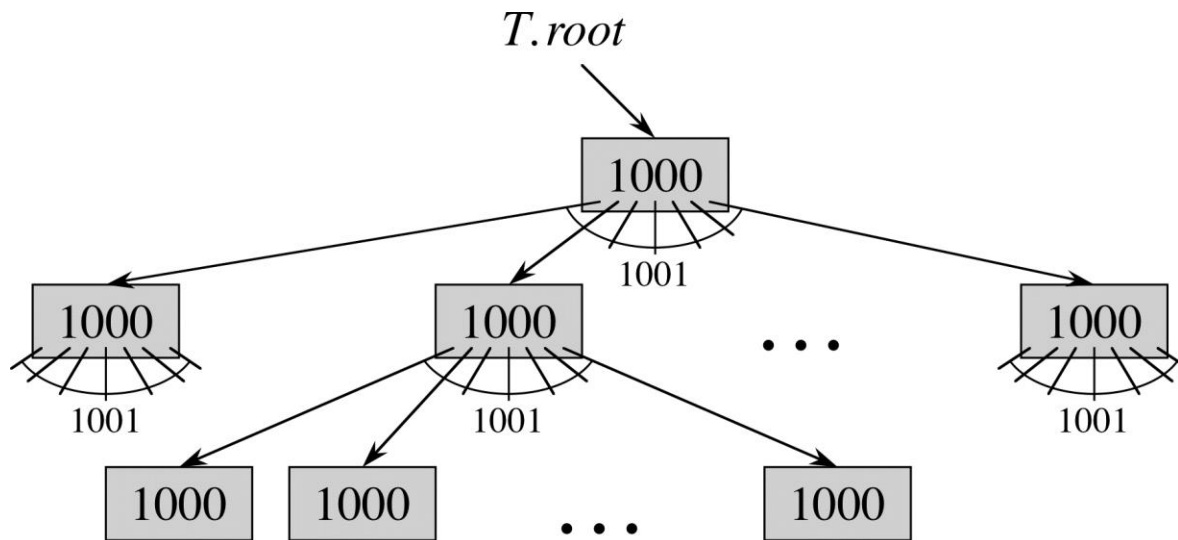
B-trees

B-trees (Bayer and McCreight, 1972)

- A **B-tree** is a **self-balancing tree** that maintains **sorted data items** and allows **basic operations in logarithmic time**.
 - Basic operations: search, sequential access, insertion, and deletion.
- It is a generalization of BST that allows for **nodes with more than two children**.
- B-tree is **commonly used in databases and file systems**.
 - It is well suited for storage systems that read and write relatively large blocks of data, such as disks.
 - Each node is filled at least 50%. In practice, it is normally ~70%.

Example: How many data items in a B-tree?

Data items in a B-tree of 1001 branches can be arranged in only 3 level
→ over one billions data items.



1 node,
1000 keys

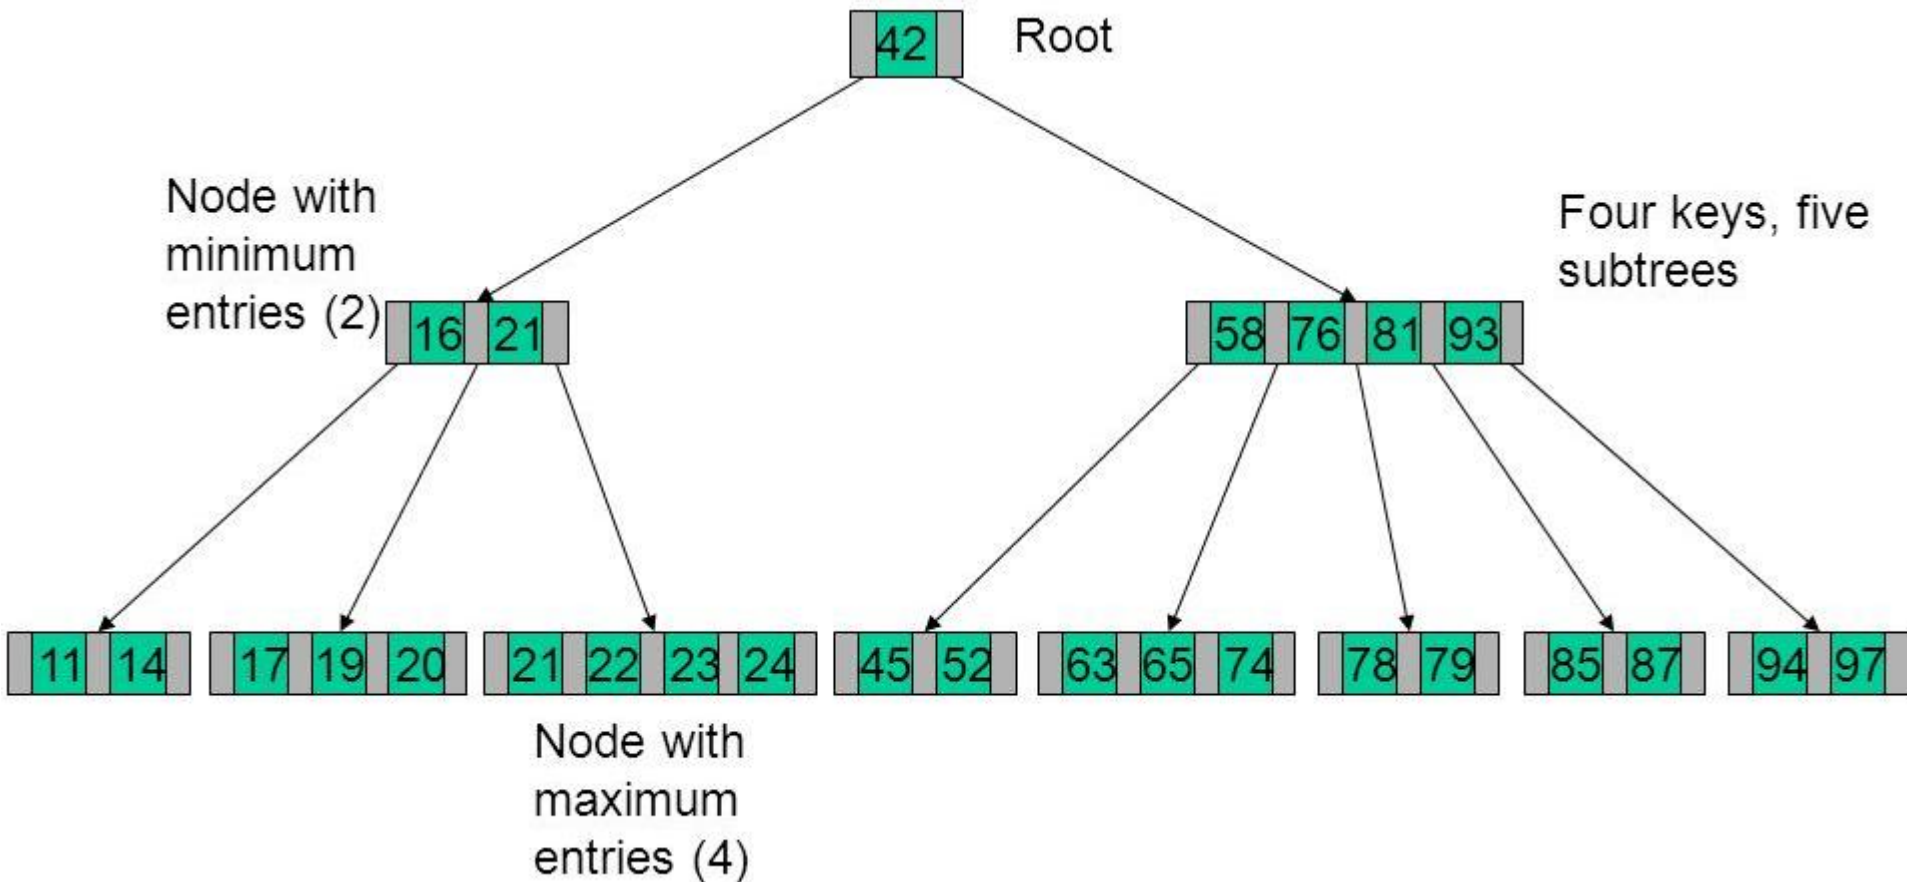
1001 nodes,
1,001,000 keys

1,002,001 nodes,
1,002,001,000 keys

B-trees: A definition

- A **B-tree of order m** ($m > 2$) is a **m -way tree** that satisfies the following conditions
 1. Every node has at most m children.
 2. Every non-leaf node (except root) has at least $\lceil m/2 \rceil$ child nodes.
 3. The root has at least two children if it is not a leaf node.
 4. A non-leaf node with k children contains $k - 1$ items.
 5. All leaves appear in the same level.
- *Note that the above B-tree is derived from Knuth's definition, while that of Bayer and McCreight is slightly different.*
- **2-3 trees** and **2-3-4 trees** are derivations of B-trees when $m = 3$ and $m = 4$, respectively.


Example: B-tree of order 5







Min # of subtrees is 3 and max is 5.

Min # of data items is 2 and max is 4.

Example: B-tree of order 4



Checkpoint 01: Which of the following is a B-tree?



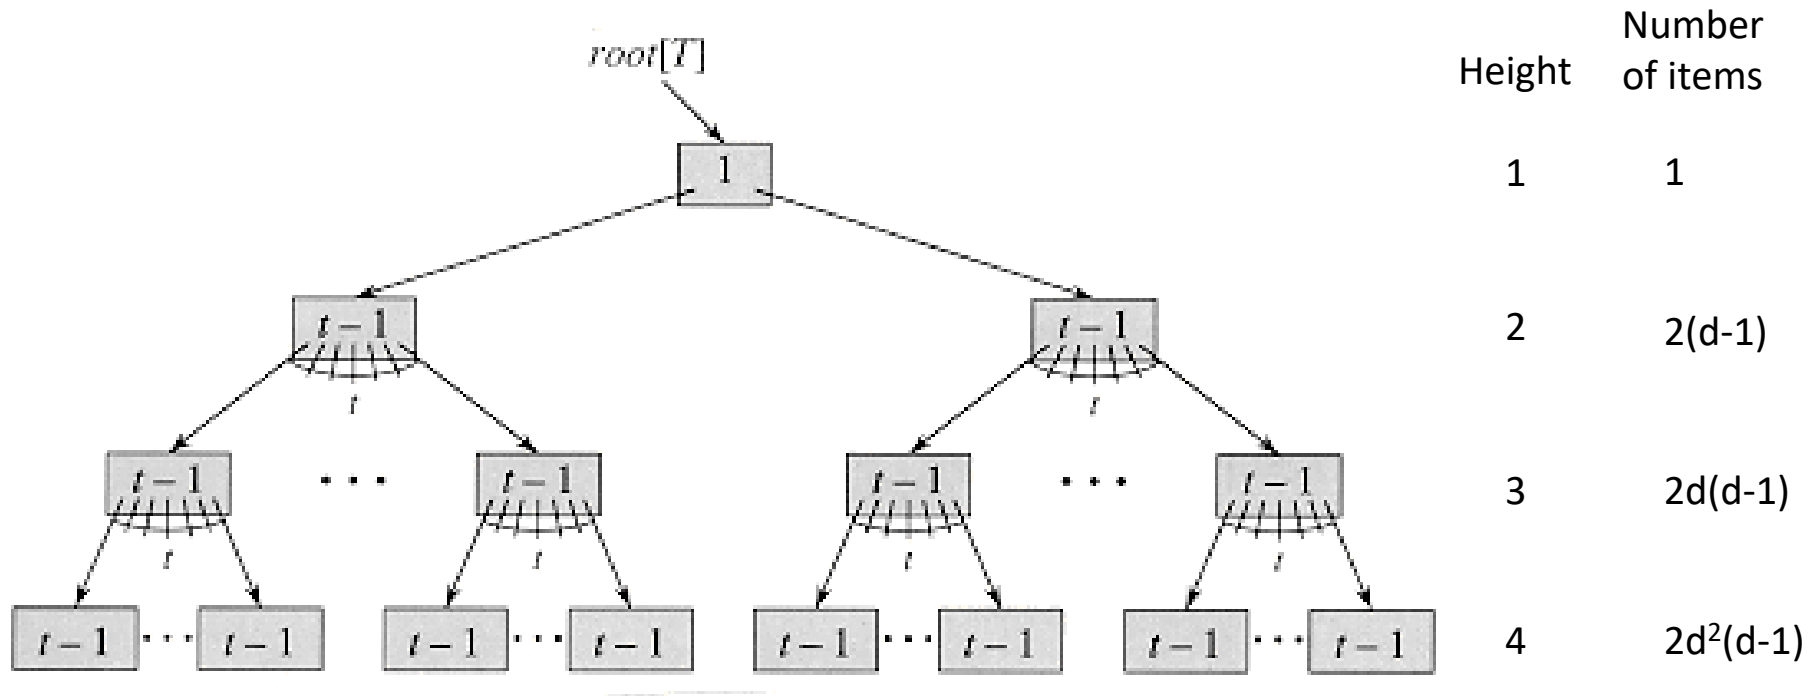
The height of a B-tree

- Let $n \geq 0$ be the number of data items in the tree.
- Let m be the maximum number of children a node can have.
 - Each node can have at most $m - 1$ data items.
- The minimum height of a B-tree is $h_{min} = \lceil \log_m(n + 1) \rceil$.
- Let d be the minimum number of children for an internal node.
For an ordinary B-tree, $d = \lceil m/2 \rceil$
- The maximum height of a B-tree is $h_{max} = \left\lceil \log_d \left(\frac{n+1}{2} \right) \right\rceil + 1$.

** Note that the height h starts from 1.*

The height of a B-tree

- In the worst case the root has only one key and two children.
- Every other node has $d - 1$ items and d children.





- The total number of items is $2d^{h-1} - 1$.

Insertion in a B-tree

- Let v be the item to be inserted into a B-tree.
- **Locate an appropriate leaf node** to insert v by traversing the tree following the order of items.
- If the **leaf node still has an empty slot**, **insert v** to this node while maintaining the order of items.
- Otherwise, **split the node**.
- The split can be back-propagated to upper nodes.
 - Worst case: the root node is split and a new root is created.




Example: Insert data item to a B-tree of order 3

Insert data item 19: The leaf node $\langle 20 \rangle$ still has an empty slot.





Example: Insert data item to a B-tree of order 3

Insert data item 21: The leaf node is full → split the node → move the middle item 20 to its parent node



Example: Insert data item to a B-tree of order 3





Insert data item 21:The parent node becomes full too → split it using the same procedure.



Checkpoint 02a: Insertion in a B-tree

What is the resultant B-tree of order 4 after inserting the keys in the following order? Insert: 5, 3, 21, 9, 1, 13, 2, 7, 10, 12, 4, 8.

Knowing that the “middle key” is right-biased.



Checkpoint 02b: Insertion in a B-tree



Insert the keys 78, 52, 81, 40, 33, 90, 85, 30 and 38 in this order in an initially empty B-tree of order 3.

Removal in a B-tree

- Let v be the item to be removed from a B-tree.
- If v is **at the leaf node**
 - If the number of items remained after removal $\geq \left\lceil \frac{m}{2} \right\rceil - 1 \rightarrow$ **Stop!**
 - Otherwise, borrow an item from a sibling node that has more than $\left\lceil \frac{m}{2} \right\rceil - 1$ items.
 - If there is no such sibling node, merge the node with its sibling node and take one data item down from the parent node.
 - If the parent node becomes insufficient, apply the same procedure.
- If v is **NOT at the leaf node**
 - Find a substation for v at some leaf node and delete the actual slot at that leaf node.



Example: Remove data item from a B-tree of order 3

Remove data item 26: The node $\langle 22, 26 \rangle$ is not a leaf node \rightarrow find a substitution.






Example: Remove data item from a B-tree of order 3

Remove data item 26: Replace 26 by 28 → delete the slot that previously contains 28.






Example: Remove data item from a B-tree of order 3

Remove data item 22: Replace 22 by 24
→ the leaf node becomes insufficient.



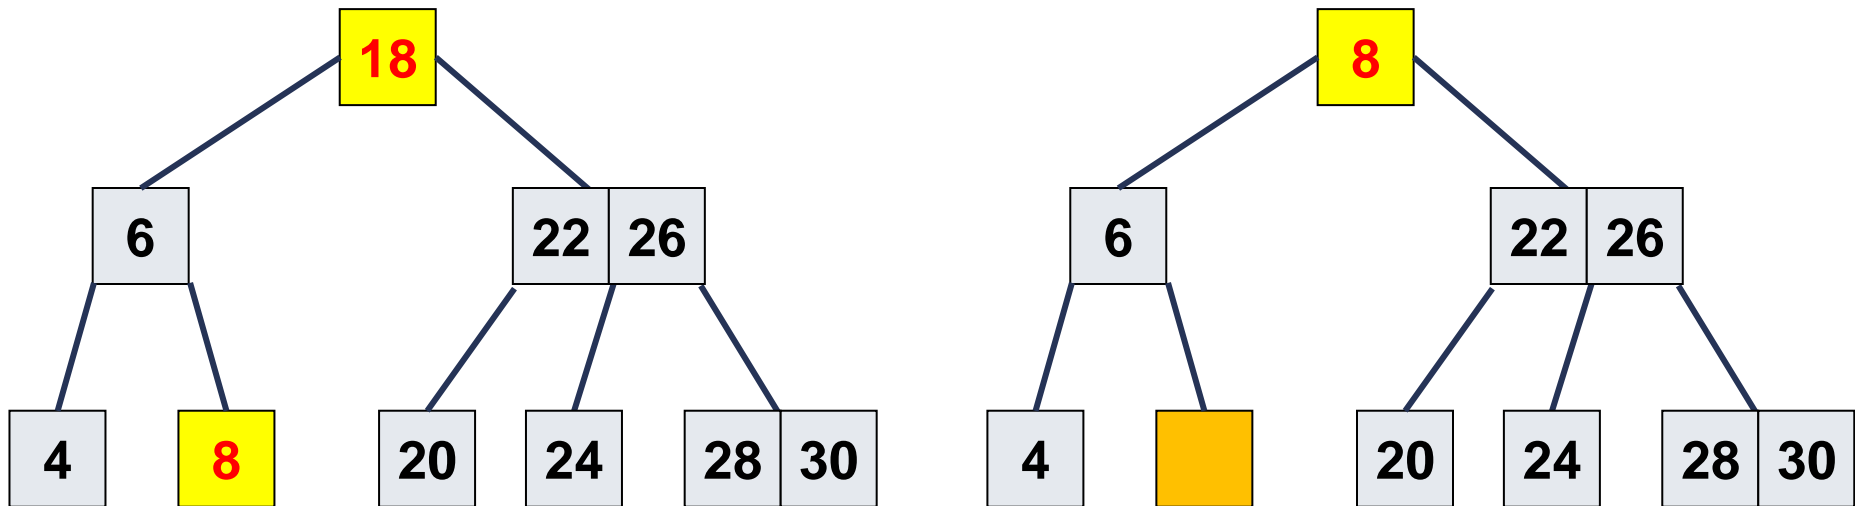
Example: Remove data item from a B-tree of order 3

Remove data item 22:borrow 28
from a sibling node




Example: Remove data item from a B-tree of order 3

Remove data item 18: Replace 18 by 8...the leaf node become insufficient.




Example: Remove data item from a B-tree of order 3

Remove data item 18: ...merge the two sibling and the data items in the parent node → the parent node becomes insufficient.



Example: Remove data item from a B-tree of order 3

Remove data item 18: ...borrow 22 in the sibling node and move corresponding subtrees.



B-trees: Implementation



```
class BNode{
    bool leaf;           // A flag for the leaf node
    unsigned int nItems;  // The number of available data items
    ItemType keys[MAX-1]; // An array of data items
    // An array of pointers to subtrees
    unsigned long pointers[MAX];
    // Redundant data to completely fill the block
    char unused[K];
}
* MAX and K are predefined constants.
```

B-tree vs. *m*-way tree

- *m*-way tree is not a balanced search tree.
 - Data insertion and removal are quite simple.
 - The tree grows toward the leaves.
- B-Tree is basically a balanced *m*-way tree.
 - Data insertion or removal may involve the split/merge of nodes.
 - It minimizes the number of accesses to external storage.
 - The tree grows towards the root.


Store a tree in an external storage

- It is essential to have an efficient search method on the external storage.
- Let t be the time to read/write a block
- $t =$ the time to move the reader to corresponding block
+ the time to read/write the block into memory



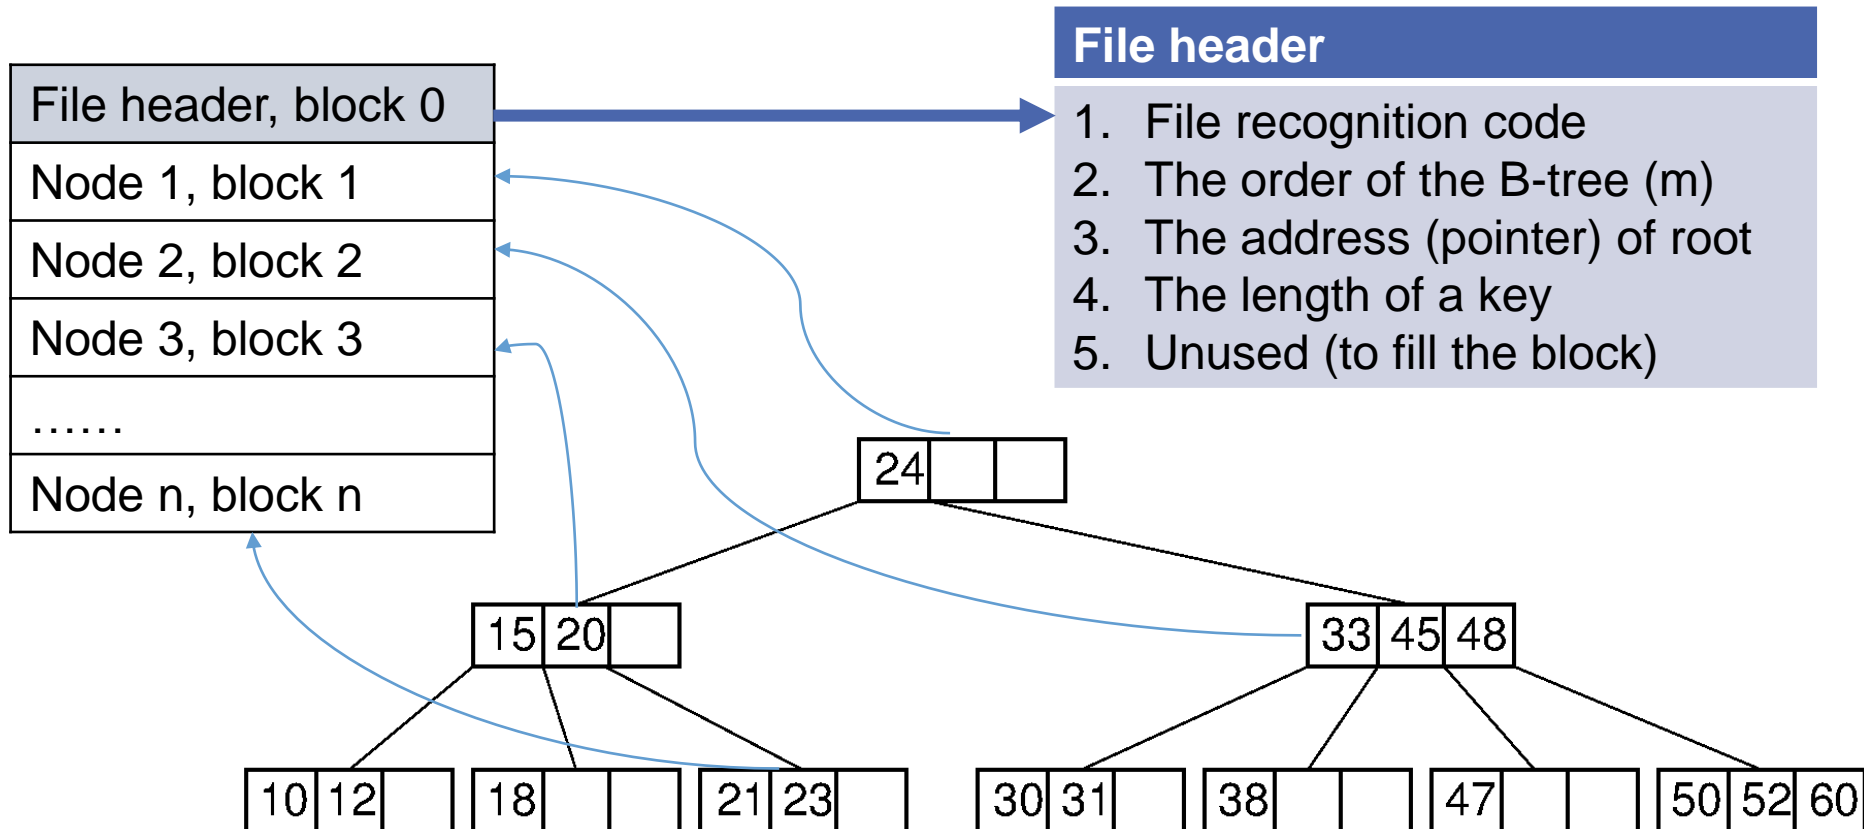
Store a tree in an external storage

- Here is a example of storing a binary tree, in which data items are grouped into blocks, to a disk storage.




Store a tree in an external storage

- For a B-tree, the root should be cached for frequent access.
 - It is not necessary to perform the READ_ROOT
 - The WRITE_ROOT is only required when the root is modified.




Store a tree in an external storage

- A B-tree with data items that have no supplementary info

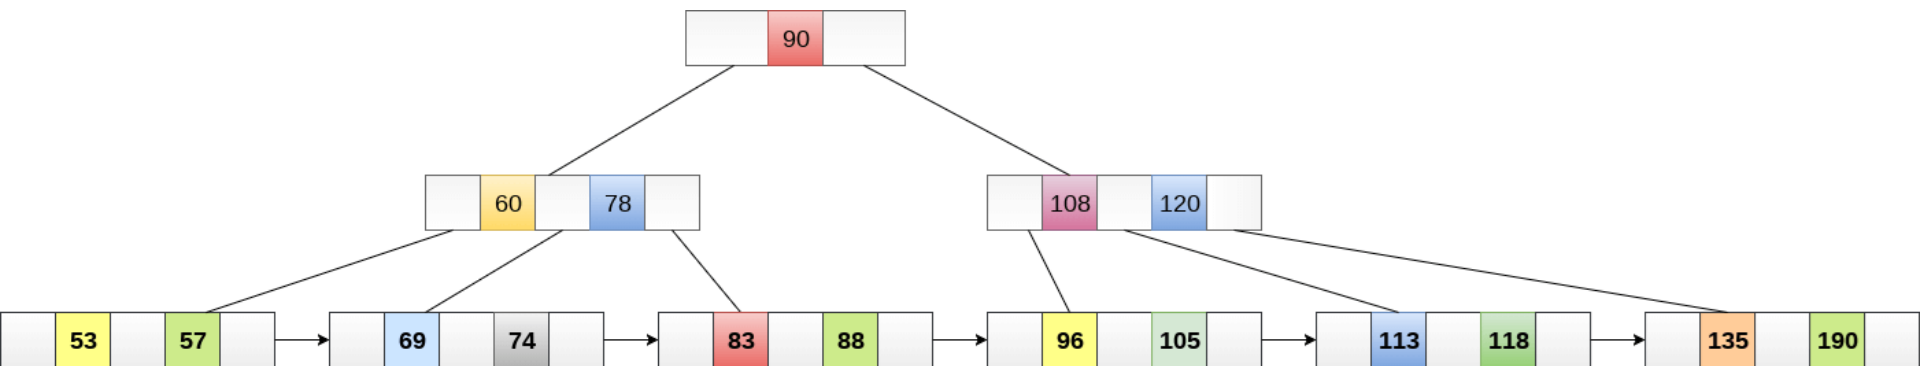


- A B-tree with data items that have certain supplementary info



B+ Trees: A variant of B-tree

- Records (data) can only be stored in the leaf nodes while internal nodes can only store the key values.
 - Meanwhile, keys and records in a B-tree both can be stored in the internal and leaf nodes.
- The leaf nodes are linked together in the form of a singly linked lists to make the search queries more efficient.
- The internal nodes (keys to access records) are stored in the main memory, and leaves are in the secondary memory.




Exercises



01. Insertion in a B-tree


- Consider the following **B-tree of order 6**.



- Draw the resulting tree after each insertion of the following keys: B, Q, L, and F

02. Removal in a B-tree

- Consider the following **B-tree of order 6**.



- Draw the resulting tree after each deletion of the following keys: F, M, G, D, and B

