

Project Report-Xv6 & Unix utilities

University of Science - VNUHCM

- **Date:** 15/02/2025
- **Subject:** *Operating System*
- **Class:** 23CLC03
- **Environment:** *Ubuntu-24.04*
- **Instructors:**
 - *Phan Quốc Kỳ*
 - *Lê Hà Minh*
 - *Lê Giang Thanh*

I. Group Information

No	Student's ID	Student's Full Name	Contribution
1	23127004	Lê Nhật Khôi	33% : <i>pingpong, primes</i>
2	23127113	Nguyễn Trần Phú Quý	33% : <i>xargs, writeups</i>
3	23127165	Nguyễn Hải Đăng	33% : <i>pingpong, find</i>

II. Submission's Details

Here are all of the assignments stated in the lab.

1. pingpong

The program creates two processes that exchange a byte over pipes.

Key system calls:

- `pipe()` : Creates pipes for bidirectional communication
- `fork()` : Creates a child process
- `read()` and `write()` : Exchange data through pipes

- `getpid()` : Gets process IDs for output

Testing

When run in the xv6 shell, the program produces the expected output:

```
$ pingpong
4: received ping
3: received pong
$
```

2. primes

We used **Eratosthenes's Sieve** to implement this assignment.

Eratosthenes's Sieve is an ancient algorithm for finding prime numbers up to any given limit.

It works by iteratively marking the multiples of each prime starting from 2, ensuring that only prime numbers remain unmarked.

The algorithm of *Eratosthenes's Sieve*

The algorithm for finding all the prime numbers to a given integer `n` is based on the following steps

1. **Step 1:** Create a list of consecutive integers from 2 to `n` and initially let `p = 2` which is the smallest prime numbers
2. **Step 2:** Remove the multiples of `p` (except `p`) such as: `2*p`, `3*p`, ...
3. **Step 3:** Find the smallest number in the list greater than `p` and set this new number as `p`
4. **Step 4:** Looping conditions:
 - If there are not any numbers larger than `p` in the list, the process stops. All of the remaining numbers in the lists are prime numbers.
 - Else, repeat from **Step 2**.

Using processes and pipes to implement *Eratosthenes's Sieve*

To implement *Eratosthenes's Sieve*, each processes is responsible for filtering a prime number. Along with processes are pipes which are responsible for passing numbers between processes. This algorithm is simulated through the following examples:

```
Input:  2  3  4  5  6  7  8  9 10 11
```

```
Process 1 (2):
```

- Print 2
- Send → 3 5 7 9 11
- Drop: 4 6 8 10

```
Process 2 (3):
```

- Print 3
- Send → 5 7 11
- Drop: 9

```
Process 3 (5):
```

- Print 5
- Send → 7 11

```
Process 4 (7):
```

- Print 7
- Send → 11

```
Process 5 (11):
```

- Print 11

3. find

On **Linux**, the basic functionality of `find` could be replicate by this command:

```
ls -R | grep "filename.ext"
```

So initially, we try combining the `ls` and `grep` commands to make `find`, but this approach failed because:

- The xv6 `ls` is missing the `-R` option and so it doesn't recursively traverse directories
- The shell's piping `|` mechanism is limited for this specific task
- This approach wouldn't handle the proper display of full pathnames

Instead, we developed a dedicated solution that correctly implements the required functionality.

Key Components

- **Directory Traversal:** Adapted code from `user/ls.c` to read directory entries
- **Recursive Search:** Implemented recursion to explore subdirectories
- **Path Handling:** Properly constructed file paths and extracted filenames
- Skipped `.` and `..` entries to avoid infinite recursion

Testing

When tested with the required directory structure, the program produces the expected output:

```
$ find . b
./b
./a/b
./a/aa/b
```

4. xargs

The program reads input lines from `stdin` and executes a specified command with each line as an additional argument.

The implementation meets the basic requirements but could be enhanced with more robust error handling and input processing.

Key Implementation Details

- Used basic system calls: `fork()`, `exec()`, `wait()`, `read()`
- Processes input one character at a time until newline
- Handles command execution by creating child processes
- Passes command-line arguments along with input line to executed command

Limitations

- Does not support the `-n` optimization found in *UNIX* `xargs`
- Maximum line length fixed at *100* characters
- No support for multiple arguments per input line
- Input buffer could overflow with very long lines

Testing

Successfully tested with basic use cases:

- `echo hello too | xargs echo bye`
- `find . b | xargs grep hello`

IV. References

- [Lab Answers](#)
- [GeeksforGeeks - Linux Commands](#)
- [File Descriptors and Open Files](#)

